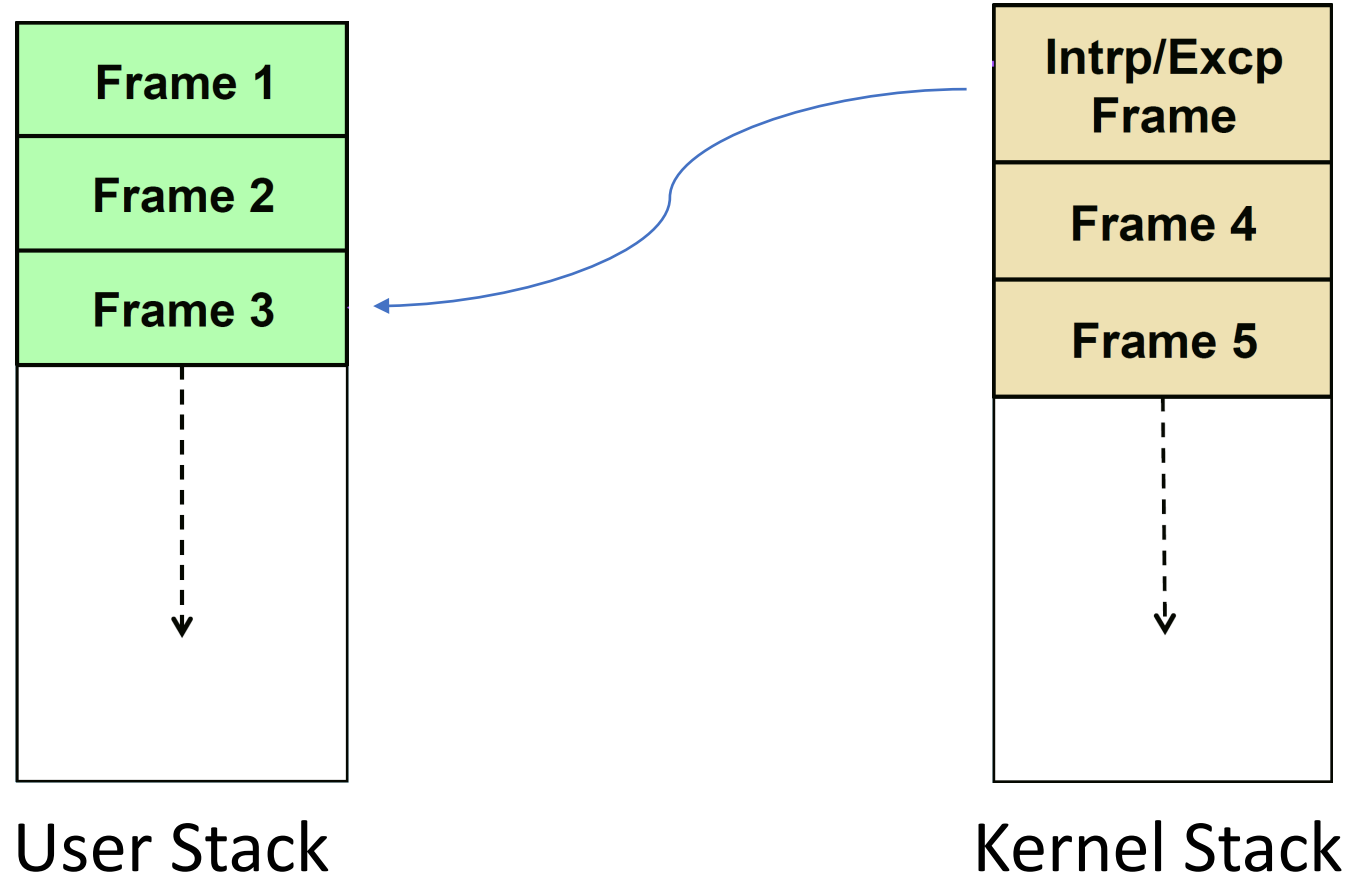


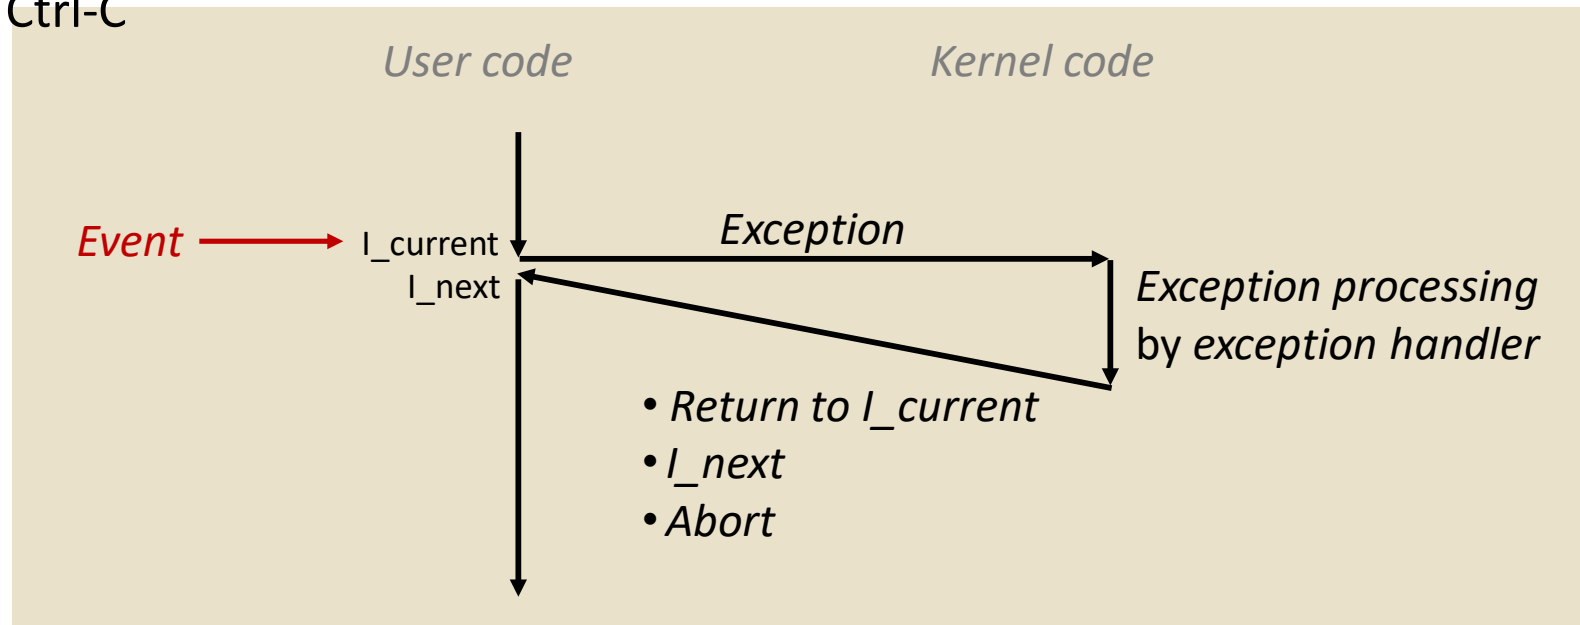
Architecture and OS, Process

One Stack Per Mode



Exceptions

- An *exception* is a transfer of control to the OS *kernel* in response to some *event*
 - Kernel is the memory-resident part of the OS
 - Examples of events: Divide by 0, arithmetic overflow, page fault, I/O request completes, typing Ctrl-C



System Calls


- Each x86-64 system call has a unique ID number
- Examples:

<i>Number</i>	<i>Name</i>	<i>Description</i>
0	read	Read file
1	write	Write file
2	open	Open file
3	close	Close file
4	stat	Get info about file
57	fork	Create process
59	execve	Execute a program
60	_exit	Terminate process
62	kill	Send signal to process

System Call Example: Opening File

- User calls: `open(filename, options)`
- Calls `__open` function, which invokes system call instruction `syscall`

```
0000000000e5d70 <__open>:  
...  
e5d79: b8 02 00 00 00    mov $0x2,%eax # open is syscall #2  
e5d7e: 0f 05             syscall      # Return value in %rax  
e5d80: 48 3d 01 f0 ff ff  cmp $0xffffffff001,%rax  
...  
e5dfa: c3               retq
```



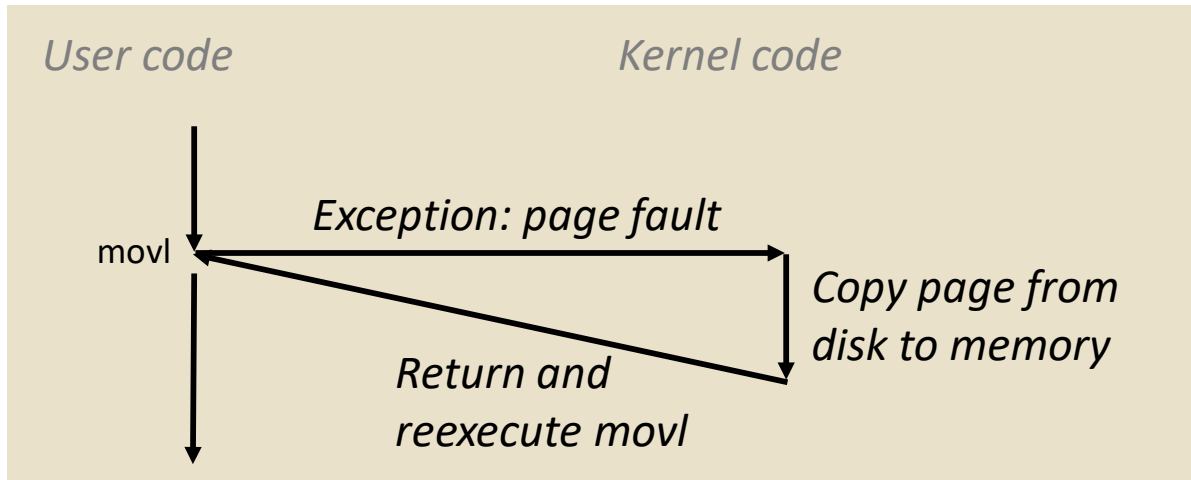
- `%rax` contains syscall number
- Other arguments in `%rdi`, `%rsi`, `%rdx`, `%r10`, `%r8`, `%r9`
- Return value in `%rax`
- Negative return value is an error : `errno`

Fault Example: Page Fault

- User writes to memory location
- That portion (page) of user's memory is currently on disk

```
int a[1000];  
main ()  
{  
    a[500] = 13;  
}
```

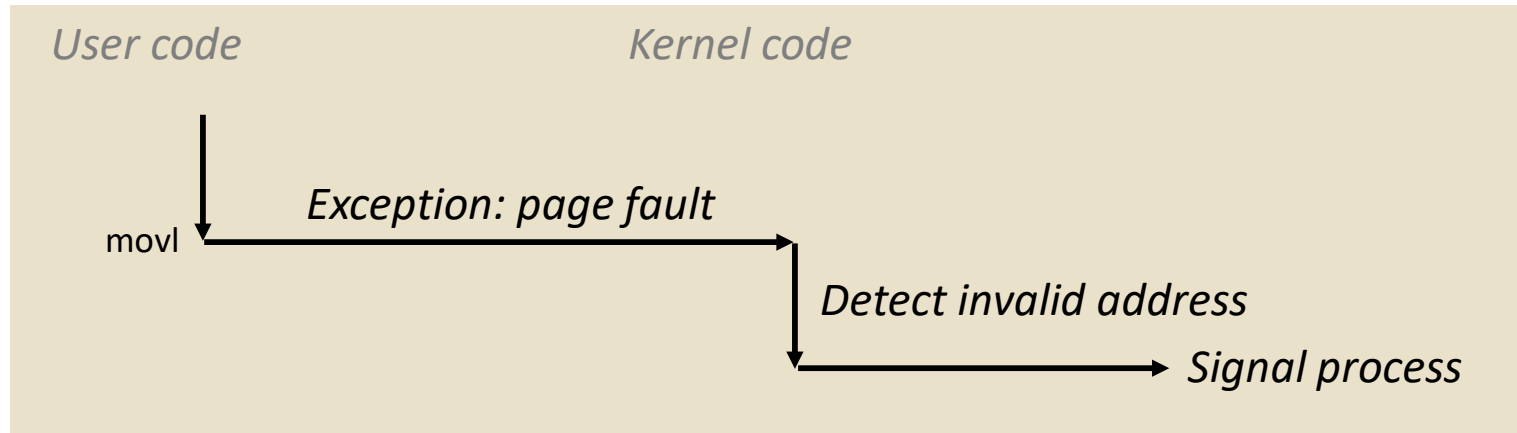
```
80483b7:      c7 05 10 9d 04 08 0d  movl    $0xd,0x8049d10
```



Fault Example: Invalid Memory Reference

```
int a[1000];  
main ()  
{  
    a[5000] = 13;  
}
```

80483b7: c7 05 60 e3 04 08 0d movl \$0xd,0x804e360



- User process exits with “segmentation fault”


Creating Your Own Processes

```
#include <unistd.h>



int main( ) {
    pid_t pid;
    if ((pid = fork()) == 0) {
        /* new process starts running here */
    }
    /* old process continues here */
}
```



Creating a Process: Before




Creating a Process: After



```
int main()
{
int value = 5;
pid_t pid;
if ((pid = fork()) == 0) { /* child process */
    value += 15;
    printf("CHILD: value = %d\n",value);
    return 0;
}
else if (pid > 0) { /* parent process */
    printf("PARENT: value = %d\n",value);
    return 0;
}
}
```

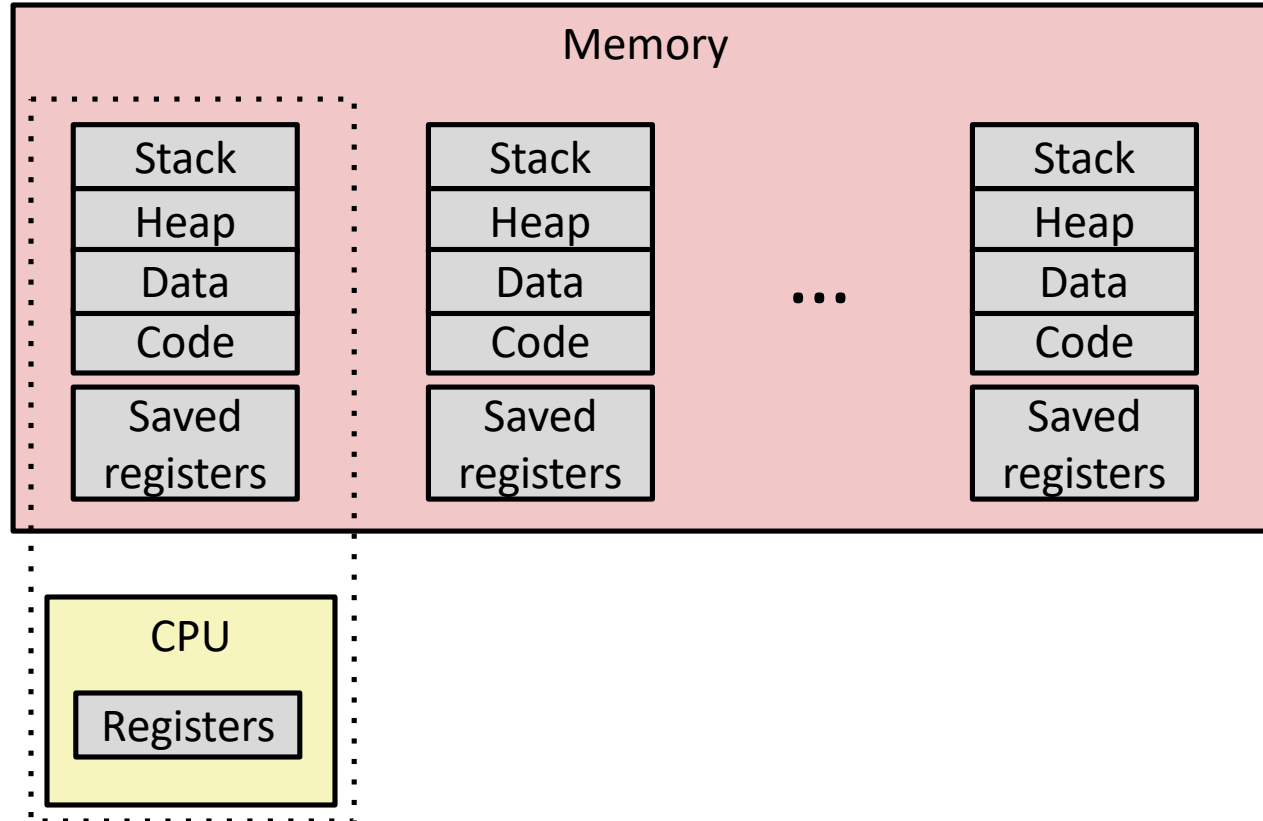
```
int main()
{
int value = 5;
pid_t pid;
if ((pid = fork()) == 0) { /* child process */
    value += 15;
    printf("CHILD: value = %d\n",value);
    return 0;
}
else if (pid > 0) { /* parent process */
    wait(NULL);
    printf("PARENT: value = %d\n",value);
    return 0;
}
}
```

Multiprocessing: The Illusion



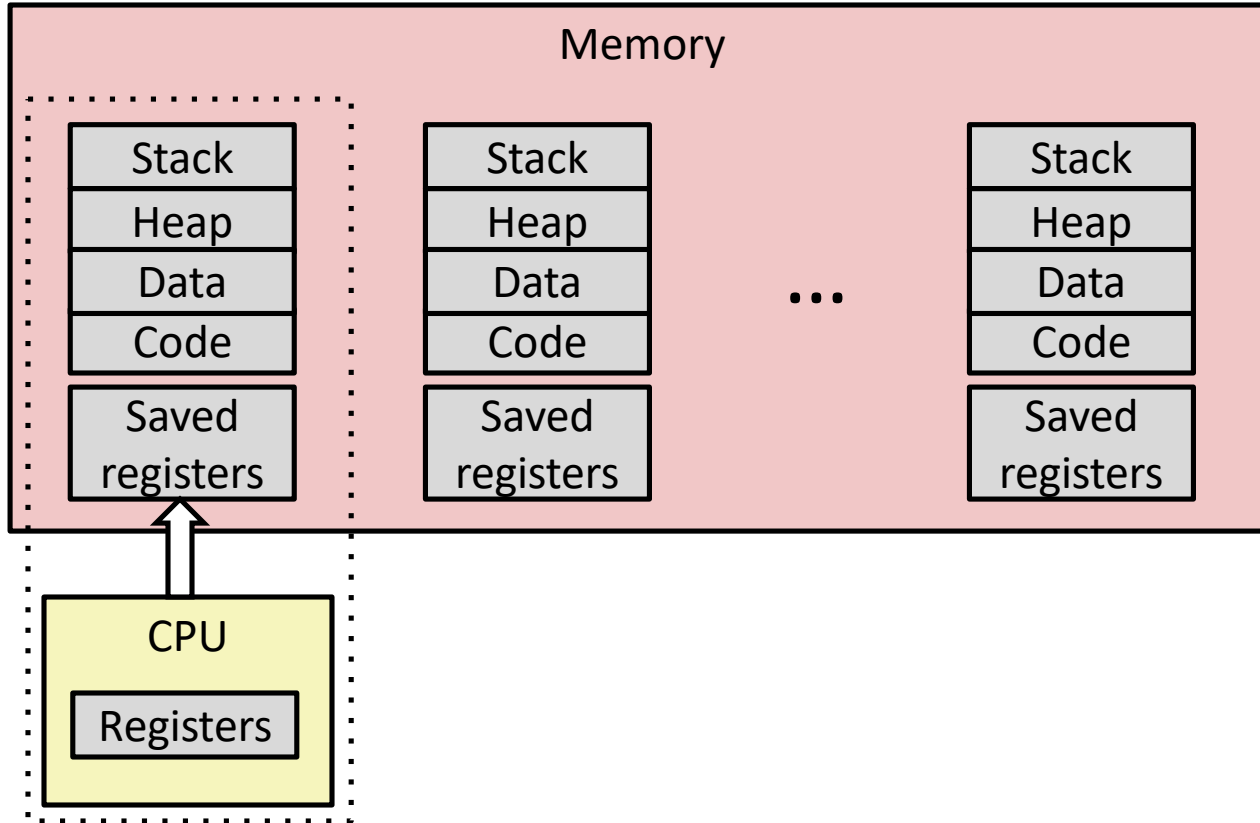
- Computer runs many processes simultaneously
 - Applications for one or more users
 - Web browsers, email clients, editors, ...

Multiprocessing: The (Traditional) Reality



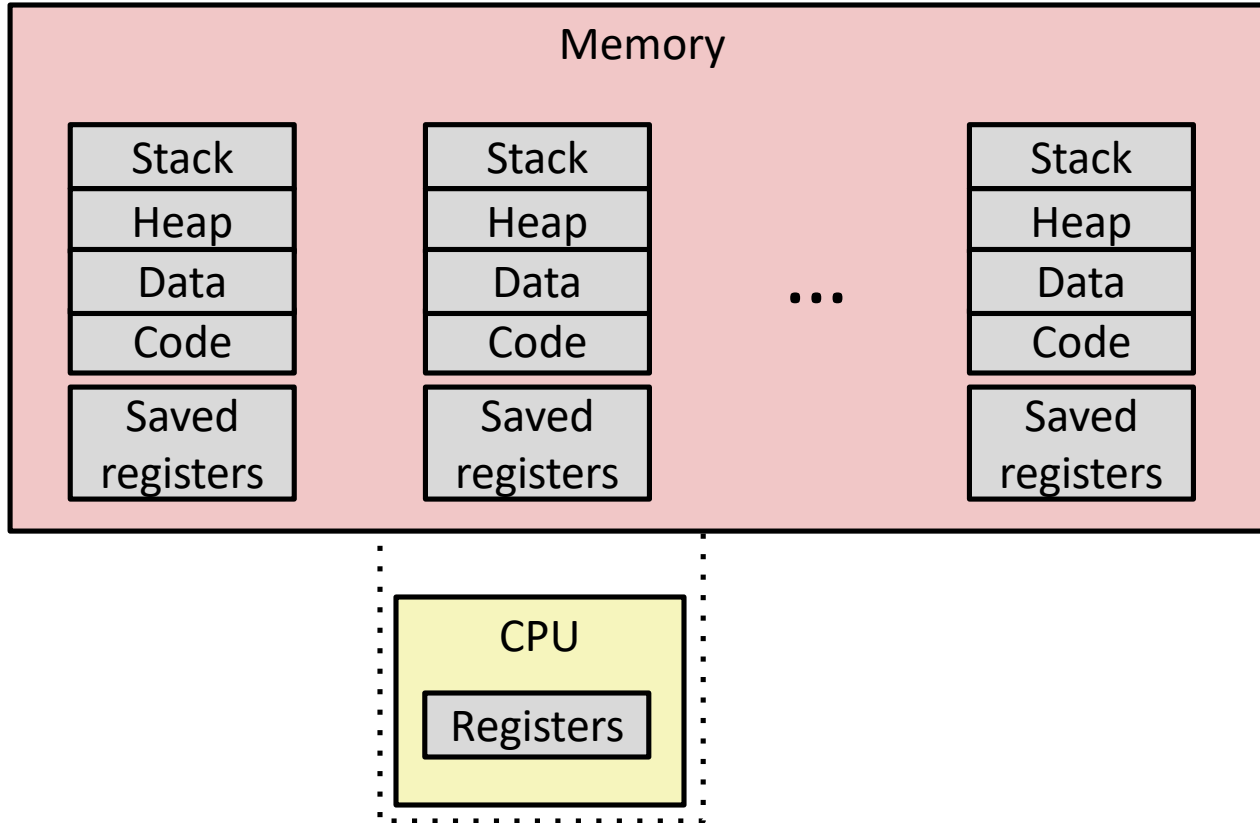
- Single processor executes multiple processes concurrently
 - Process executions interleaved (multitasking)
 - Address spaces managed by virtual memory system (later in course)
 - Register values for nonexecuting processes saved in memory

Multiprocessing: The (Traditional) Reality



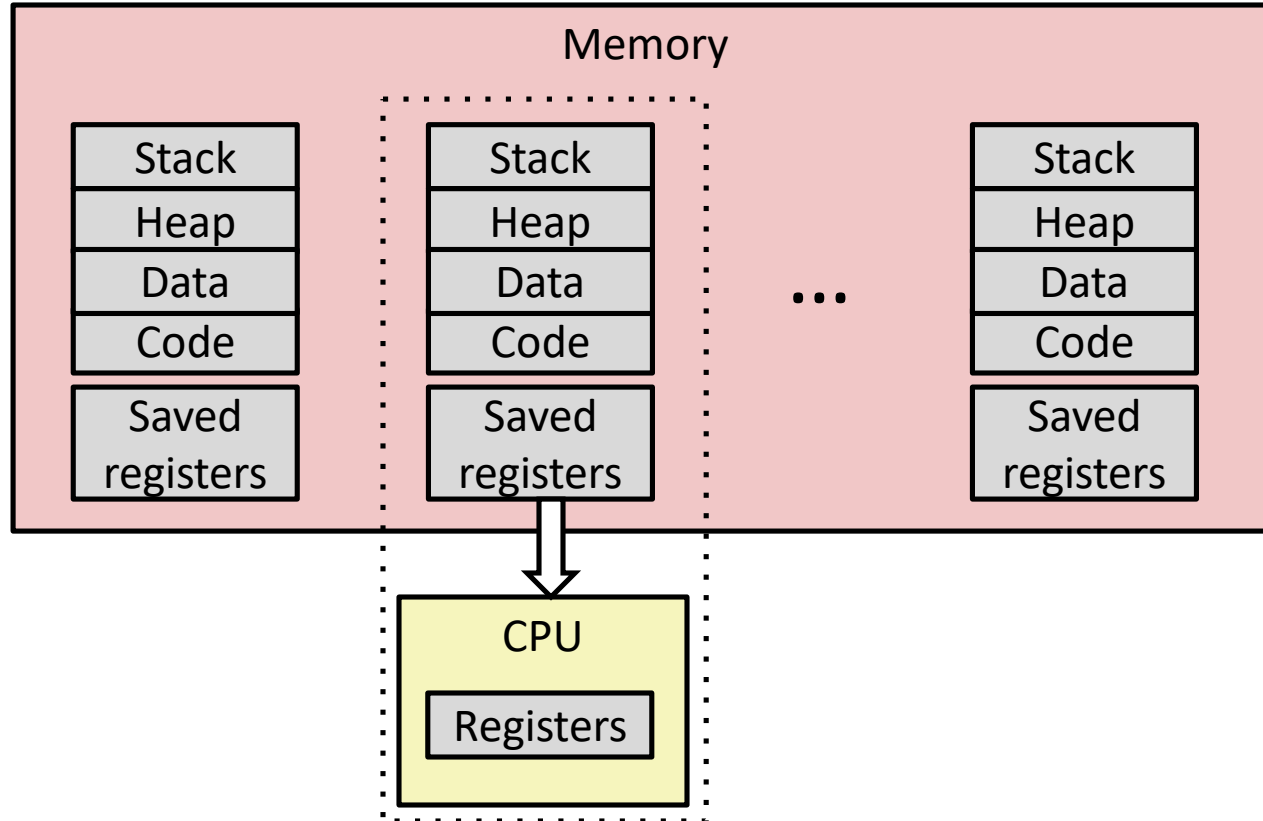
- Save current registers in memory

Multiprocessing: The (Traditional) Reality




- Schedule next process for execution

Multiprocessing: The (Traditional) Reality



- Load saved registers and switch address space (context switch)

Multiprocessing: The (Modern) Reality



- Multicore processors
 - Multiple CPUs on single chip
 - Share main memory (and some caches)
 - Each can execute a separate process
 - Scheduling of processors onto cores done by kernel