

Process

* Some materials adapted from CS:APP, Prof. Bryant, O'Hallaron, Doeppner, Galvin, etc.

Processes

- **A program in execution**
- **Containers for programs**
 - **virtual memory**
 - address space
 - **scheduling**
 - one or more threads of control
 - **file references**
 - open files
 - **and lots more!**

Obtaining Process IDs

- `pid_t getpid(void)`
 - Returns PID of current process
- `pid_t getppid(void)`
 - Returns PID of parent process

Creating Your Own Processes

```
#include <unistd.h>

int main( ) {
    pid_t pid;
    if ((pid = fork()) == 0) {
        /* new process starts running here */
    }
    /* old process continues here */
}
```



System Calls

- **Sole direct interface between user and kernel**
- **Implemented as library routines that execute *trap***
- **instructions to enter kernel**
- **Errors indicated by returns of -1 ; error code is in global variable *errno***

```
if (write(fd, buffer, bufsize) == -1) {  
    // error!  
    printf("error %d\n", errno);  
}
```

fork Example: Two consecutive forks

```
void fork2 ()
{
    printf("L0\n");
    fork();
    printf("L1\n");
    fork();
    printf("Bye\n");
}
```

forks.c

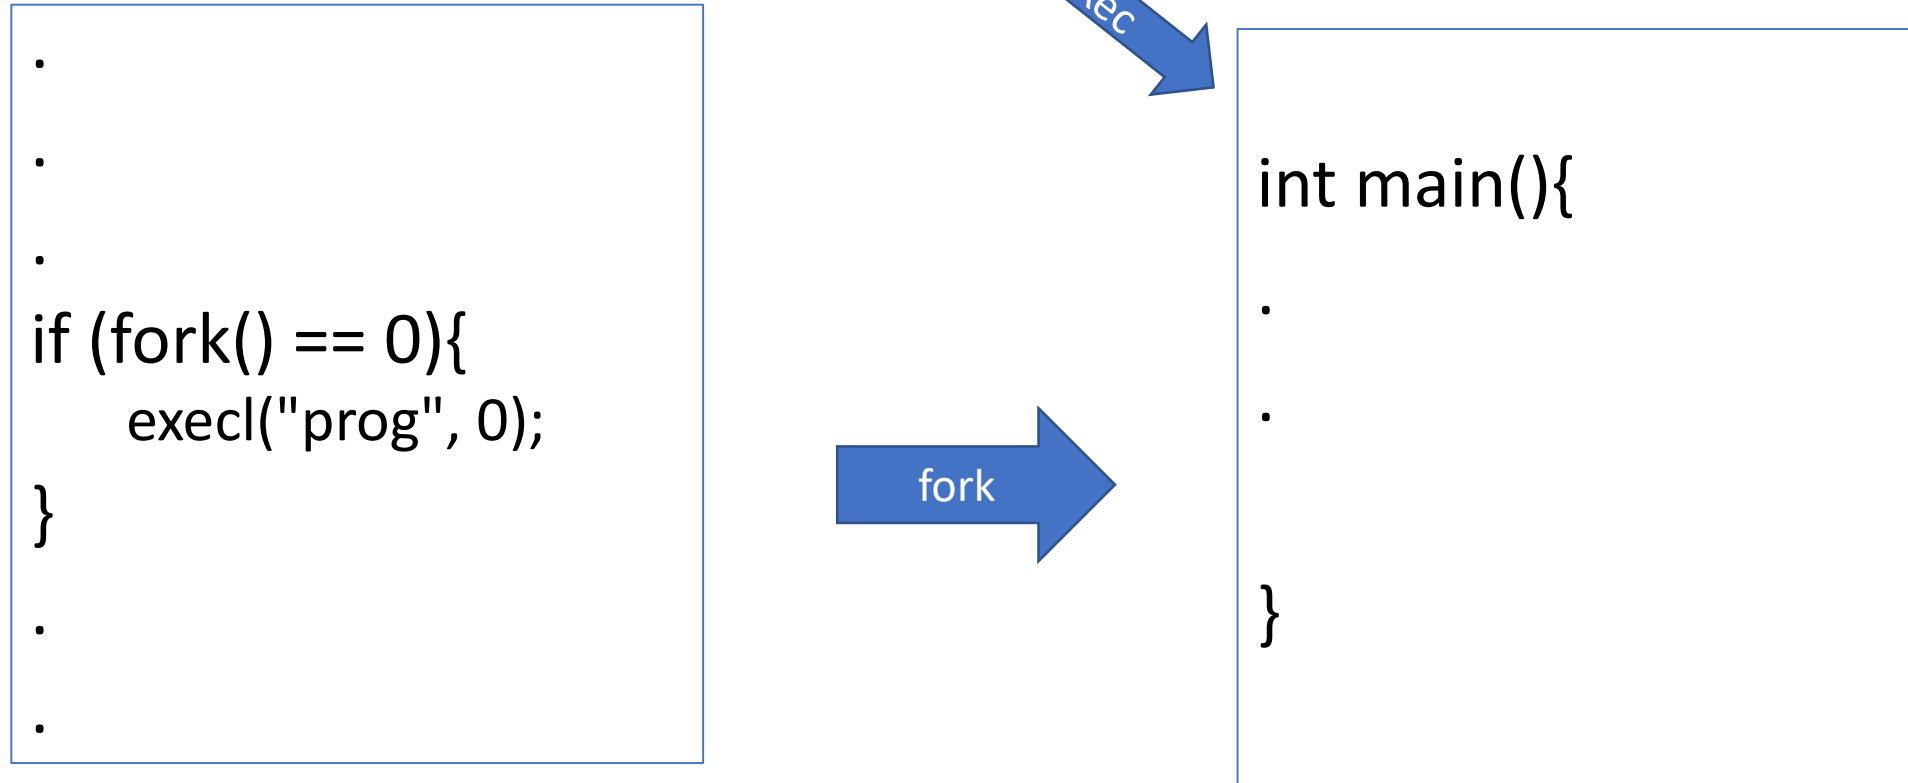
Feasible output:

L0
L1
Bye
Bye
L1
Bye
Bye

Infeasible output:

L0
Bye
L1
Bye
L1
Bye
Bye

Putting Programs into Processes




Execv

- Family of related routines
 - we concentrate on one:
 - `execv(program, argv)`

```
char *argv[] = {“./MyProg”, “12”, (void *)0};  
if (fork() == 0) {  
    execv("./MyProg", argv);  
}
```


Loading a New Image



A Random Program ...

```
int main(int argc, char *argv[]) {  
    int i;  
    int stop = atoi(argv[1]);  
    for (i = 0; i < stop; i++)  
        printf("%d\n", rand());  
    return 0;  
}
```

Passing It Arguments

- From the shell

```
$ ./random 12
```

- **From a C program**

```
if (fork() == 0) {  
    char *argv[] = {"./random", "12", (void *)0};  
    execv("./random", argv);  
}
```


```
if (fork() == 0) {  
    char *argv[] = {"/random", "12", (void *)0};  
    execv("/random", argv);  
    printf("random done\n");  
}
```

Question: Will the *printf* statement be executed, after random completes?

Receiving Arguments

```
int main(int argc, char *argv[]) {  
    int i;  
    int stop = atoi(argv[1]);  
    for (i = 0; i < stop; i++)  
        printf("%d\n", rand());  
    return 0;  
}
```

argv



Shell

- **How does the shell invoke your program?**

```
if (fork() == 0) {  
    char *argv = {"/./random", "12", (void *)0};  
    execv("/./random", argv);  
}  
/* what does the shell do here??? */
```

Wait

```
#include <unistd.h>
#include <sys/wait.h>
...
pid_t pid;
int status;
...
if ((pid = fork()) == 0) {
    char *argv[] = {"/random", "12", (void *)0};
    execv("/random", argv);
}
waitpid(pid, &status, 0);
```

Shell: To Wait or Not To Wait ...

- \$ who

```
    if ((pid = fork()) == 0) {  
        char *argv[] = {"who", 0};  
        execv("who", argv);  
    }  
    waitpid(pid, &status, 0);
```

...

- \$ who &

```
    if ((pid = fork()) == 0) {  
        char *argv[] = {"who", 0};  
        execv("who", argv);  
    }
```


Zombie Example

```
void fork7() {  
    if (fork() == 0) {  
        /* Child */  
        printf("Terminating Child, PID = %d\n", getpid());  
        exit(0);  
    } else {  
        printf("Running Parent, PID = %d\n", getpid());  
        while (1)  
            ; /* Infinite loop */  
    }  
}
```

```
linux> ./fork7 &  
Running Parent, PID = 6639  
Terminating Child, PID = 6640  
linux> ps  
  PID TTY          TIME CMD  
 6585 ttyp9      00:00:00 tcsh  
 6639 ttyp9      00:00:03 forks  
 6640 ttyp9      00:00:00 forks <defunct>  
 6641 ttyp9      00:00:00 ps  
linux> kill 6639  
[1]      Terminated  
linux> ps  
  PID TTY          TIME CMD  
 6585 ttyp9      00:00:00 tcsh  
 6642 ttyp9      00:00:00 ps
```

- ps shows child process as "defunct" (i.e., a zombie)

- Killing parent allows child to be reaped by init

Reaping Child Processes

- Idea
 - When process terminates, it still consumes system resources
 - Examples: Exit status, various OS tables
 - Called a “zombie”
 - Living corpse, half alive and half dead
- Reaping
 - Performed by parent on terminated child (using `wait` or `waitpid`)
 - Parent is given exit status information
 - Kernel then deletes zombie child process
- What if parent doesn't reap?
 - If any parent terminates without reaping a child, then the orphaned child will be reaped by `init` process (`pid == 1`)
 - So, only need explicit reaping in long-running processes
 - e.g., shells and servers

Non-terminating Child Example

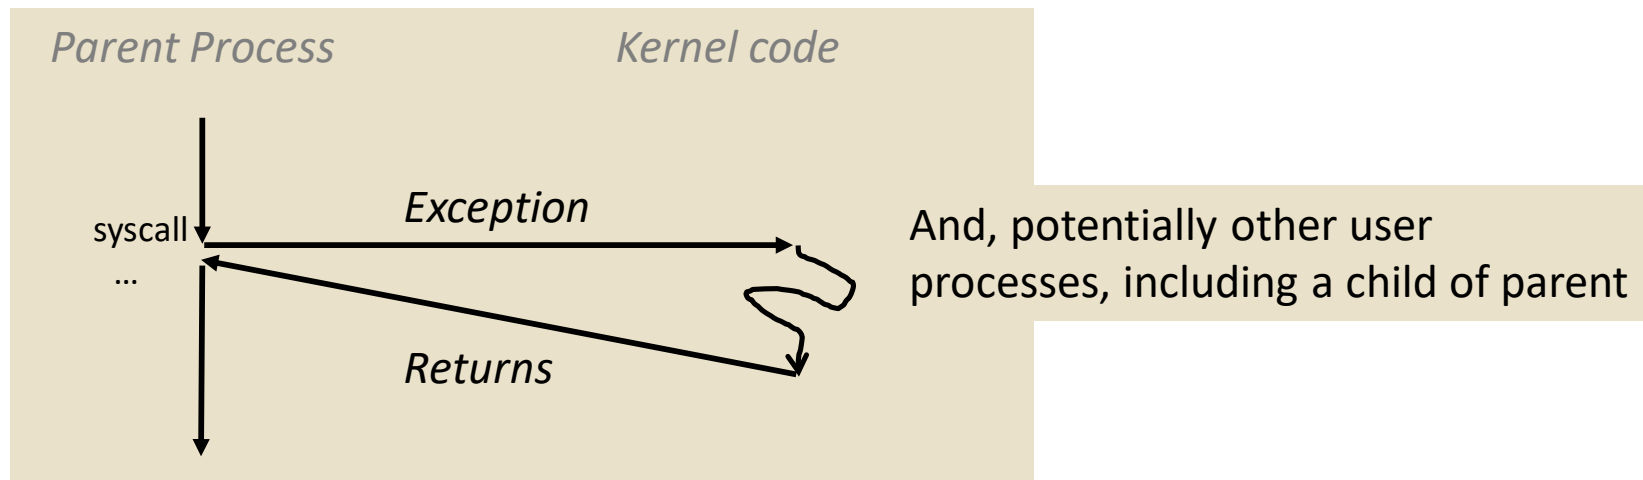
```
void fork8()
{
    if (fork() == 0) {
        /* Child */
        printf("Running Child, PID = %d\n",
               getpid());
        while (1)
            ; /* Infinite loop */
    } else {
        printf("Terminating Parent, PID = %d\n",
               getpid());
        exit(0);
    }
}
```

```
linux> ./forks8
Terminating Parent, PID = 6675
Running Child, PID = 6676
linux> ps
  PID TTY          TIME CMD
 6585 tttyp9      00:00:00 tcsh
 6676 tttyp9      00:00:06 forks
 6677 tttyp9      00:00:00 ps
linux> kill 6676
linux> ps
  PID TTY          TIME CMD
 6585 tttyp9      00:00:00 tcsh
 6678 tttyp9      00:00:00 ps
```

- Child process still active even though parent has terminated
- Must kill child explicitly, or else will keep running indefinitely

`wait`: Synchronizing with Children

- Parent reaps a child by calling the `wait` function
- `int wait(int *child_status)`
 - Suspends current process until one of its children terminates



`wait`: Synchronizing with Children

- Parent reaps a child by calling the `wait` function
- `int wait(int *child_status)`
 - Suspends current process until one of its children terminates
 - Return value is the **pid** of the child process that terminated
 - If **`child_status != NULL`**, then the integer it points to will be set to a value that indicates reason the child terminated and the exit status:
 - Checked using macros defined in `wait.h`
 - `WIFEXITED`, `WEXITSTATUS`, `WIFSIGNALED`, `WTERMSIG`, `WIFSTOPPED`, `WSTOPSIG`, `WIFCONTINUED`

wait: Synchronizing with Children

```
void fork9() {  
    int child_status;  
  
    if (fork() == 0) {  
        printf("HC: hello from child\n");  
        exit(0);  
    } else {  
        printf("HP: hello from parent\n");  
        wait(&child_status);  
        printf("CT: child has terminated\n");  
    }  
    printf("Bye\n");  
}
```

forks.c

Feasible:

HC	HP
HP	HC
CT	CT
Bye	Bye

Infeasible:

HP
CT
Bye
HC

Exit

```
#include <unistd.h>
#include <stdlib.h>
#include <sys/wait.h>
int main( ) {
    pid_t pid;
    int status;
    if ((pid = fork()) == 0) {
        if (do_work() == 1)
            exit(0); /* success! */
        else
            exit(1); /* failure ... */
    }
    waitpid(pid, &status, 0);
    /* WEXITSTATUS(status) extracts it */
}
```

Summary

- Exceptions
 - Events that require nonstandard control flow
 - Generated externally (interrupts) or internally (traps and faults)
- Processes
 - At any given time, system has multiple active processes
 - Only one can execute at a time on any single core
 - Each process appears to have total control of processor + private memory space

Summary (cont.)

- Spawning processes
 - Call `fork`
 - One call, two returns
- Process completion
 - Call `exit`
 - One call, no return
- Reaping and waiting for processes
 - Call `wait` or `waitpid`
- Loading and running programs
 - Call `execve` (or variant)
 - One call, (normally) no return