



Memory Management



Memory Technologies


| | Capacity | Latency | Cost/GB |
|-----------|---------------|---------|----------|
| Registers | 1000s of bits | 20 ps | \$\$\$\$ |
| SRAM | ~10KB-10MB | 1-10 ns | ~\$1000 |
| DRAM | ~10GB | 80 us | ~10 |
| Hard disk | ~1TB | 10 ms | ~0.10 |

Different technologies have vastly different tradeoffs.

small, low latency

large, high latency

Memory Hierarchy Levels





Levels in a typical memory hierarchy.

Closer to CPU, memory becomes faster , smaller, as well as more expensive per bit

Locality

- **Principle of Locality:** Programs tend to use data and instructions with addresses near or equal to those they have used recently
- **Temporal locality:**
 - Recently referenced items are likely to be referenced again in the near future
- **Spatial locality:**
 - Items with nearby addresses tend to be referenced close together in time



Locality Example

```
sum = 0;  
for (i = 0; i < n; i++)  
    sum += a[i];  
return sum;
```

- Data references

- Reference array elements in succession
- Reference variable **sum** each iteration.

Spatial locality

Temporal locality

- Instruction references

- Reference instructions in sequence.
- Cycle through loop repeatedly.

Spatial locality

Temporal locality

Locality Example

- **Claim:** Being able to look at code and get a qualitative sense of its locality is a key skill for a professional programmer.
- **Question:** Which function has good locality with respect to array *a*?


```
int sum_array_rows(int a[M][N])
{
    int i, j, sum = 0;

    for (i = 0; i < M; i++)
        for (j = 0; j < N; j++)
            sum += a[i][j];
    return sum;
}
```

```
int sum_array_cols(int a[M][N])
{
    int i, j, sum = 0;


    for (j = 0; j < N; j++)
        for (i = 0; i < M; i++)
            sum += a[i][j];
    return sum;
}
```

Memory Hierarchy Levels




- Block (aka line): unit of copying
 - May be multiple words
- If accessed data is present in upper level
 - Hit: access satisfied by upper level
 - Hit ratio: hits/accesses

Memory Hierarchy Levels



- If accessed data is absent
 - Miss:
 - Miss ratio: $\text{misses} / \text{accesses}$
 $= 1 - \text{hit ratio}$
 - Copy the block from lower level to upper level
 - Then accessed data is available in upper level

Memory Hierarchy Levels



- If accessed data is absent
 - Miss
 - Miss ratio: $\text{misses/accesses} = 1 - \text{hit ratio}$
 - Copy the block from lower level to upper level
 - Then accessed data is available in upper level

Cache Performance Metrics

■ Hit Time

- Time to deliver a line in the cache to the processor
 - includes time to determine whether the line is in the cache
- Typical numbers:
 - 1~4 clock cycle for L1
 - 10~100 clock cycles for L2

■ Miss Penalty

- Additional time required because of a miss
 - typically 50-200 cycles for main memory

Let's think about those numbers

- Huge difference between a hit and a miss
 - Could be 100x, if just L1 and main memory
- Would you believe 99% hits is twice as good as 97%?
 - Consider:
 - cache hit time of 1 cycle
 - miss penalty of 100 cycles
 - Average access time: $\text{HitTime} + \text{MissRatio} \times \text{MissPenalty}$
 - 97% hits: 1 cycle + 0.03 x 100 cycles = **4 cycles**
 - 99% hits: 1 cycle + 0.01 x 100 cycles = **2 cycles**
- This is why “miss rate” is used instead of “hit rate”

Cache Memory

- Cache memory
 - The level of the memory hierarchy closest to the CPU
- Given accesses X_1, \dots, X_{n-1}, X_n

| |
|-----------|
| X_4 |
| X_1 |
| X_{n-2} |
| |
| X_{n-1} |
| X_2 |
| |
| X_3 |

a. Before the reference to X_n

| |
|-----------|
| X_4 |
| X_1 |
| X_{n-2} |
| |
| X_{n-1} |
| X_2 |
| X_n |
| X_3 |

b. After the reference to X_n

- How do we know if the data is present?
- Where do we look?



Direct Mapped Cache

Memory

00001

00101

01001

01101

10001

10101

11001

11101

Cache

000

001

010

011

100

101

110

111

- Location determined by address
- Direct mapped: only one choice
 - $(\text{Block address in memory}) \bmod (\text{\#Blocks in cache})$
- Use low-order address bits

Tags and Valid Bits

- How do we know which particular block is stored in a cache location?
 - Store block address as well as the data
 - Actually, only need the high-order bits
 - Called the tag
- What if there is no data in a location?
 - Valid bit: 1 = present, 0 = not present
 - Initially: valid bit = 0;

Cache Example

- 8-blocks, 1 word/block, direct mapped
- Initial state

| Index | V | Tag | Data |
|-------|---|-----|------|
| 000 | 0 | | |
| 001 | 0 | | |
| 010 | 0 | | |
| 011 | 0 | | |
| 100 | 0 | | |
| 101 | 0 | | |
| 110 | 0 | | |
| 111 | 0 | | |

Cache Example

Cache: word address means block address

| Index | V | Tag | Data |
|-------|---|-----|------|
| 000 | 0 | | |
| 001 | 0 | | |
| 010 | 0 | | |
| 011 | 0 | | |
| 100 | 0 | | |
| 101 | 0 | | |
| 110 | 0 | | |
| 111 | 0 | | |

Memory:

| Word addr | Binary addr | Hit/miss | Cache index |
|-----------|-------------|----------|-------------|
| 22 | 10 110 | Miss | 110 |

Cache Example

Cache: word address means block address

| Index | V | Tag | Data |
|-------|---|-----|------------|
| 000 | 0 | | |
| 001 | 0 | | |
| 010 | 0 | | |
| 011 | 0 | | |
| 100 | 0 | | |
| 101 | 0 | | |
| 110 | 1 | 10 | Mem[10110] |
| 111 | 0 | | |

Memory:

| Word addr | Binary addr | Hit/miss | Cache index |
|-----------|-------------|----------|-------------|
| 22 | 10 110 | Miss | 110 |

Cache Example

Cache: word address means block address

| Index | V | Tag | Data |
|-------|---|-----|------------|
| 000 | 0 | | |
| 001 | 0 | | |
| 010 | 0 | | |
| 011 | 0 | | |
| 100 | 0 | | |
| 101 | 0 | | |
| 110 | 1 | 10 | Mem[10110] |
| 111 | 0 | | |

Memory:

| Word addr | Binary addr | Hit/miss | Cache index |
|-----------|-------------|----------|-------------|
| 26 | 11 010 | Miss | 010 |

Cache Example

Cache: word address means block address

| Index | V | Tag | Data |
|-------|---|-----|------------|
| 000 | 0 | | |
| 001 | 0 | | |
| 010 | 1 | 11 | Mem[11010] |
| 011 | 0 | | |
| 100 | 0 | | |
| 101 | 0 | | |
| 110 | 1 | 10 | Mem[10110] |
| 111 | 0 | | |

Memory:

| Word addr | Binary addr | Hit/miss | Cache index |
|-----------|-------------|----------|-------------|
| 26 | 11 010 | Miss | 010 |

Cache Example

Cache: word address means block address

| Index | V | Tag | Data |
|-------|---|-----|------------|
| 000 | 0 | | |
| 001 | 0 | | |
| 010 | 1 | 11 | Mem[11010] |
| 011 | 0 | | |
| 100 | 0 | | |
| 101 | 0 | | |
| 110 | 1 | 10 | Mem[10110] |
| 111 | 0 | | |

Memory:

| Word addr | Binary addr | Hit/miss | Cache index |
|-----------|-------------|----------|-------------|
| 22 | 10 110 | Hit | 110 |
| 26 | 11 010 | Hit | 010 |

Cache Example

Cache: word address means block address

| Index | V | Tag | Data |
|-------|---|-----|------------|
| 000 | 0 | | |
| 001 | 0 | | |
| 010 | 1 | 11 | Mem[11010] |
| 011 | 0 | | |
| 100 | 0 | | |
| 101 | 0 | | |
| 110 | 1 | 10 | Mem[10110] |
| 111 | 0 | | |

Memory:

| Word addr | Binary addr | Hit/miss | Cache index |
|-----------|-------------|----------|-------------|
| 18 | 10 010 | Miss | 010 |

Cache Example

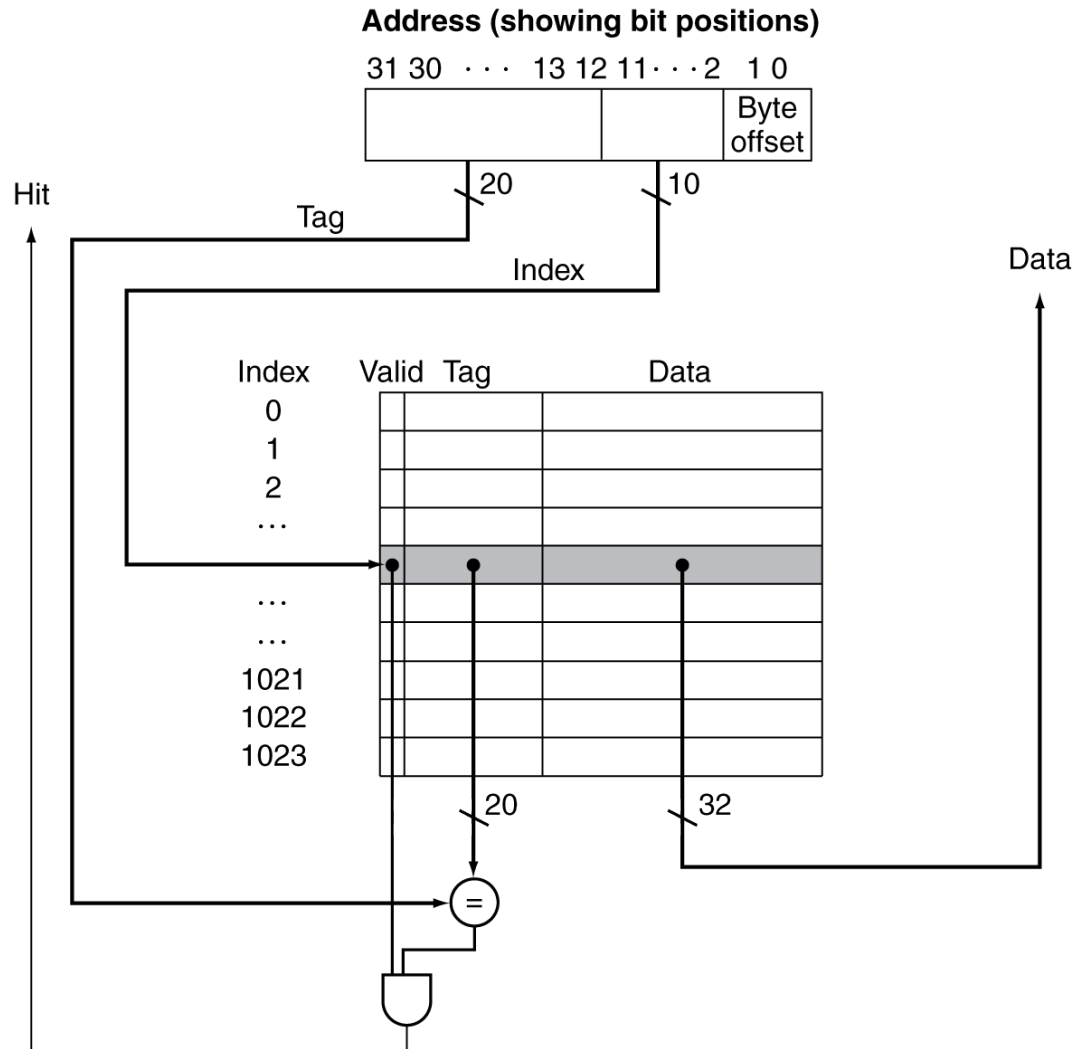
Cache: word address means block address

| Index | V | Tag | Data |
|------------|----------|-----------|-------------------|
| 000 | 0 | | |
| 001 | 0 | | |
| 010 | 1 | 10 | Mem[10010] |
| 011 | 0 | | |
| 100 | 0 | | |
| 101 | 0 | | |
| 110 | 1 | 10 | Mem[10110] |
| 111 | 0 | | |

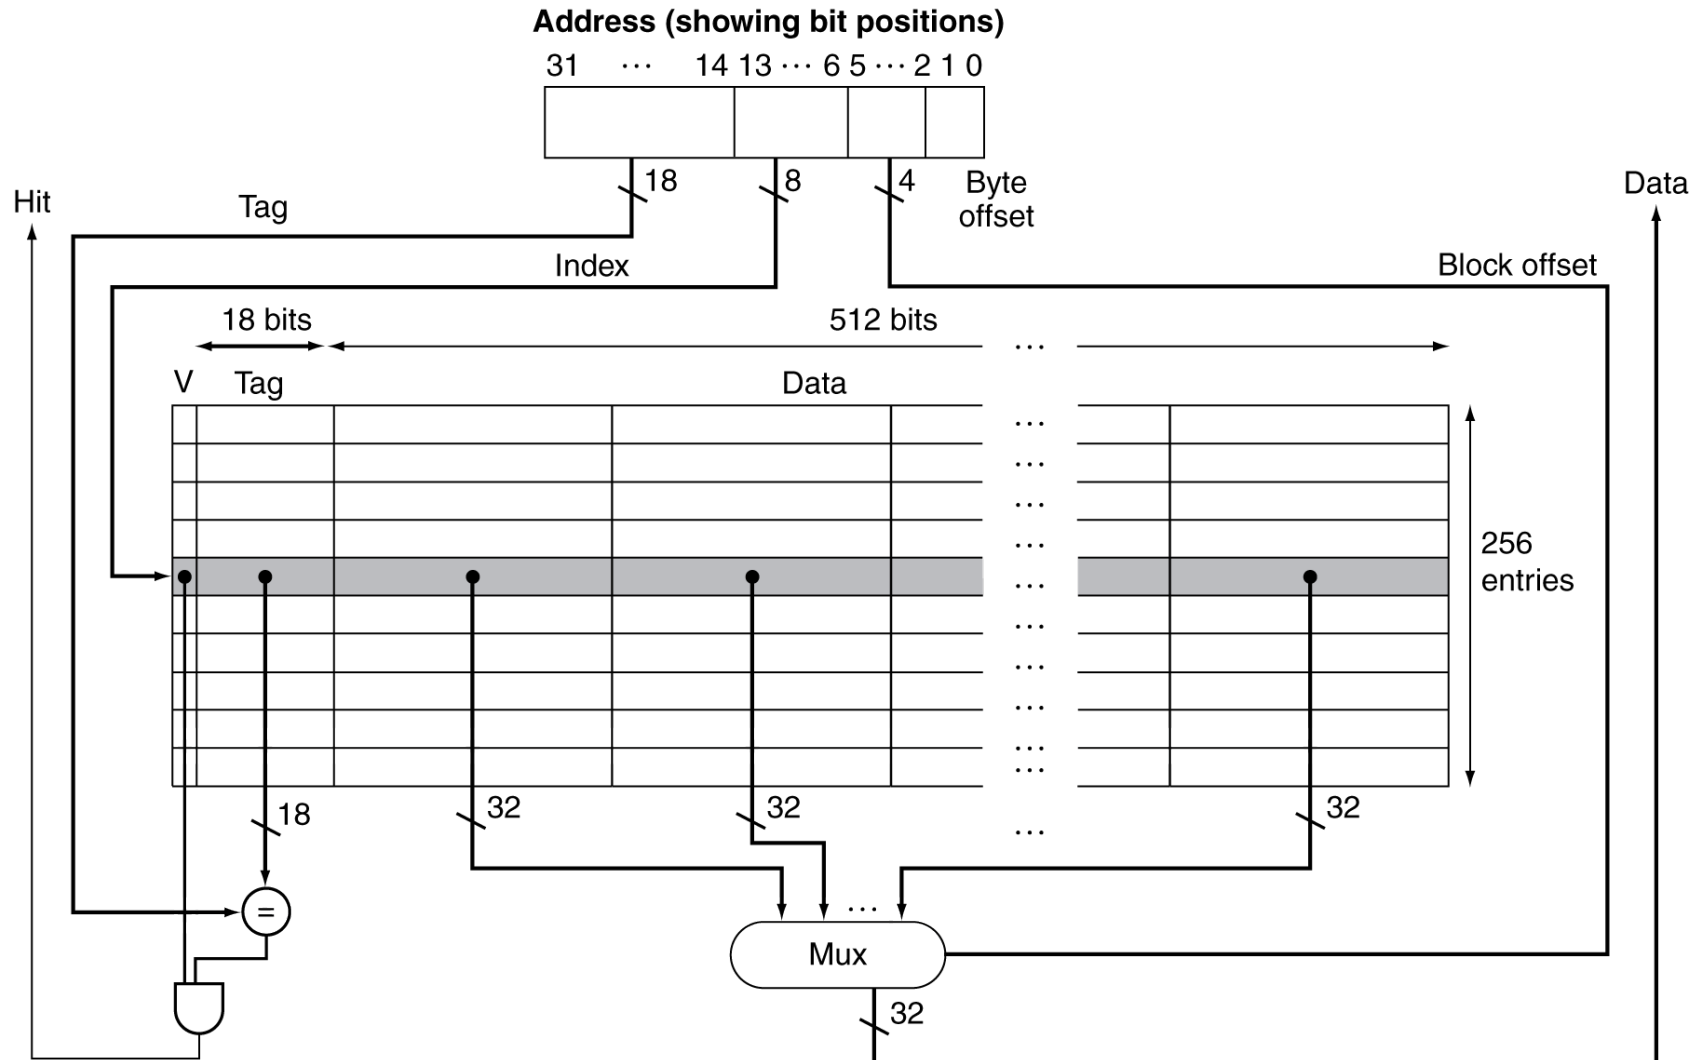
Memory:

| Word addr | Binary addr | Hit/miss | Cache index |
|-----------|-------------|----------|-------------|
| 18 | 10 010 | Miss | 010 |

Address Subdivision



Example: Intrinsic FastMATH




Block Size Tradeoffs


- Larger block sizes
 - Take advantage of spatial locality
 - Incur larger miss penalty since it takes longer to transfer the block into the cache
 - Can increase the average hit time and miss rate

Associative Cache Example


Direct mapped



Set associative



Fully associative



Reducing Cache Miss: #1 Associative Caches

- Direct mapped cache
 - A memory block maps to exactly one cache block
- Fully associative
 - Allow a given block to go in any cache entry
 - Requires all entries to be searched at once
 - Comparator per entry (expensive)
- n -way set associative
 - Each set contains n entries
 - Block address determines which set
 - (Block address) modulo (#Sets in cache)
 - Search all entries in a given set at once
 - n comparators (less expensive)

Spectrum of Associativity

- For a cache with 8 entries

**One-way set associative
(direct mapped)**

| Block | Tag | Data |
|-------|-----|------|
| 0 | | |
| 1 | | |
| 2 | | |
| 3 | | |
| 4 | | |
| 5 | | |
| 6 | | |
| 7 | | |

Two-way set associative

| Set | Tag | Data | Tag | Data |
|-----|-----|------|-----|------|
| 0 | | | | |
| 1 | | | | |
| 2 | | | | |
| 3 | | | | |

Four-way set associative

| Set | Tag | Data | Tag | Data | Tag | Data | Tag | Data |
|-----|-----|------|-----|------|-----|------|-----|------|
| 0 | | | | | | | | |
| 1 | | | | | | | | |

Eight-way set associative (fully associative)

| Tag | Data | Tag | Data | Tag | Data | Tag | Data | Tag | Data | Tag | Data | Tag | Data | Tag | Data |
|-----|------|-----|------|-----|------|-----|------|-----|------|-----|------|-----|------|-----|------|
| | | | | | | | | | | | | | | | |

Associativity Example

- Compare 4-block caches
 - Direct mapped, 2-way set associative, fully associative
 - Block access sequence: 0, 8, 0, 6, 8
- Direct mapped

| Block address | Cache index | Hit/miss | Cache content after access | | | |
|---------------|-------------|----------|----------------------------|---|---|---|
| | | | 0 | 1 | 2 | 3 |
| 0 | | | | | | |
| 8 | | | | | | |
| 0 | | | | | | |
| 6 | | | | | | |
| 8 | | | | | | |

Associativity Example

- Compare 4-block caches
 - Direct mapped, 2-way set associative, fully associative
 - Block access sequence: 0, 8, 0, 6, 8
- Direct mapped

| Block address | Cache index | Hit/miss | Cache content after access | | | |
|---------------|-------------|----------|----------------------------|---|---|---|
| | | | 0 | 1 | 2 | 3 |
| 0 | 0 | | | | | |
| 8 | 0 | | | | | |
| 0 | 0 | | | | | |
| 6 | 2 | | | | | |
| 8 | 0 | | | | | |

Associativity Example

- Compare 4-block caches
 - Direct mapped, 2-way set associative, fully associative
 - Block access sequence: 0, 8, 0, 6, 8
- Direct mapped

| Block address | Cache index | Hit/miss | Cache content after access | | | |
|---------------|-------------|----------|----------------------------|---|---|---|
| | | | 0 | 1 | 2 | 3 |
| 0 | 0 | miss | Mem[0] | | | |
| 8 | 0 | miss | Mem[8] | | | |
| 0 | 0 | | | | | |
| 6 | 2 | | | | | |
| 8 | 0 | | | | | |

Associativity Example

- Compare 4-block caches
 - Direct mapped, 2-way set associative, fully associative
 - Block access sequence: 0, 8, 0, 6, 8
- Direct mapped

| Block address | Cache index | Hit/miss | Cache content after access | | | |
|---------------|-------------|----------|----------------------------|---|---|---|
| | | | 0 | 1 | 2 | 3 |
| 0 | 0 | miss | Mem[0] | | | |
| 8 | 0 | miss | Mem[8] | | | |
| 0 | 0 | miss | Mem[0] | | | |
| 6 | 2 | | | | | |
| 8 | 0 | | | | | |

Associativity Example

- Compare 4-block caches
 - Direct mapped, 2-way set associative, fully associative
 - Block access sequence: 0, 8, 0, 6, 8
- Direct mapped

| Block address | Cache index | Hit/miss | Cache content after access | | | |
|---------------|-------------|----------|----------------------------|---|--------|---|
| | | | 0 | 1 | 2 | 3 |
| 0 | 0 | miss | Mem[0] | | | |
| 8 | 0 | miss | Mem[8] | | | |
| 0 | 0 | miss | Mem[0] | | | |
| 6 | 2 | miss | Mem[0] | | Mem[6] | |
| 8 | 0 | | | | | |

Associativity Example

- Compare 4-block caches
 - Direct mapped, 2-way set associative, fully associative
 - Block address access sequence: 0, 8, 0, 6, 8
- Direct mapped

| address | Cache index | Hit/miss | Cache content after access | | | |
|---------|-------------|----------|----------------------------|---|--------|---|
| | | | 0 | 1 | 2 | 3 |
| 0 | 0 | miss | Mem[0] | | | |
| 8 | 0 | miss | Mem[8] | | | |
| 0 | 0 | miss | Mem[0] | | | |
| 6 | 2 | miss | Mem[0] | | Mem[6] | |
| 8 | 0 | miss | Mem[8] | | Mem[6] | |

Associativity Example

- 2-way set associative

| address | Cache index | Hit/miss | Cache content after access | | | |
|---------|-------------|----------|----------------------------|--|-------|--|
| | | | Set 0 | | Set 1 | |
| 0 | | | | | | |
| 8 | | | | | | |
| 0 | | | | | | |
| 6 | | | | | | |
| 8 | | | | | | |

Associativity Example

- 2-way set associative

| Block address | Cache index | Hit/miss | Cache content after access | | | |
|---------------|-------------|----------|----------------------------|--|-------|--|
| | | | Set 0 | | Set 1 | |
| 0 | 0 | | | | | |
| 8 | 0 | | | | | |
| 0 | 0 | | | | | |
| 6 | 0 | | | | | |
| 8 | 0 | | | | | |

Associativity Example

- 2-way set associative

| Block address | Cache index | Hit/miss | Cache content after access | | | |
|---------------|-------------|----------|----------------------------|--|-------|--|
| | | | Set 0 | | Set 1 | |
| 0 | 0 | miss | Mem[0] | | | |
| 8 | 0 | | | | | |
| 0 | 0 | | | | | |
| 6 | 0 | | | | | |
| 8 | 0 | | | | | |

Associativity Example

- 2-way set associative

| Block address | Cache index | Hit/miss | Cache content after access | | | |
|---------------|-------------|----------|----------------------------|--------|-------|--|
| | | | Set 0 | | Set 1 | |
| 0 | 0 | miss | Mem[0] | | | |
| 8 | 0 | miss | Mem[0] | Mem[8] | | |
| 0 | 0 | | | | | |
| 6 | 0 | | | | | |
| 8 | 0 | | | | | |

Associativity Example

- 2-way set associative

| Block address | Cache index | Hit/miss | Cache content after access | | | |
|---------------|-------------|----------|----------------------------|--------|-------|--|
| | | | Set 0 | | Set 1 | |
| 0 | 0 | miss | Mem[0] | | | |
| 8 | 0 | miss | Mem[0] | Mem[8] | | |
| 0 | 0 | hit | Mem[0] | Mem[8] | | |
| 6 | 0 | | | | | |
| 8 | 0 | | | | | |

Associativity Example

- 2-way set associative

| Block address | Cache index | Hit/miss | Cache content after access | | | |
|---------------|-------------|----------|----------------------------|--------|-------|--|
| | | | Set 0 | | Set 1 | |
| 0 | 0 | miss | Mem[0] | | | |
| 8 | 0 | miss | Mem[0] | Mem[8] | | |
| 0 | 0 | hit | Mem[0] | Mem[8] | | |
| 6 | 0 | miss | Mem[0] | Mem[6] | | |
| 8 | 0 | miss | Mem[8] | Mem[6] | | |

Associativity Example

- Fully associative

| Block address | | Hit/miss | Cache content after access | | | |
|---------------|--|----------|----------------------------|--|--|--|
| 0 | | | | | | |
| 8 | | | | | | |
| 0 | | | | | | |
| 6 | | | | | | |
| 8 | | | | | | |

Associativity Example

- Fully associative

| Block address | | Hit/miss | Cache content after access | | | |
|---------------|--|----------|----------------------------|--|--|--|
| 0 | | miss | Mem[0] | | | |
| 8 | | | | | | |
| 0 | | | | | | |
| 6 | | | | | | |
| 8 | | | | | | |

Associativity Example

- Fully associative

| Block address | | Hit/miss | Cache content after access | | | |
|---------------|--|----------|----------------------------|--------|--|--|
| 0 | | miss | Mem[0] | | | |
| 8 | | miss | Mem[0] | Mem[8] | | |
| 0 | | | | | | |
| 6 | | | | | | |
| 8 | | | | | | |

Associativity Example

- Fully associative

| Block address | | Hit/miss | Cache content after access | | | |
|---------------|--|----------|----------------------------|--------|--|--|
| 0 | | miss | Mem[0] | | | |
| 8 | | miss | Mem[0] | Mem[8] | | |
| 0 | | hit | Mem[0] | Mem[8] | | |
| 6 | | | | | | |
| 8 | | | | | | |

Associativity Example

- Fully associative

| Block address | | Hit/miss | Cache content after access | | | |
|---------------|--|----------|----------------------------|--------|--------|--|
| 0 | | miss | Mem[0] | | | |
| 8 | | miss | Mem[0] | Mem[8] | | |
| 0 | | hit | Mem[0] | Mem[8] | | |
| 6 | | miss | Mem[0] | Mem[8] | Mem[6] | |
| 8 | | | | | | |

Associativity Example


- Fully associative

| Block address | | Hit/miss | Cache content after access | | | |
|---------------|--|----------|----------------------------|--------|--------|--|
| 0 | | miss | Mem[0] | | | |
| 8 | | miss | Mem[0] | Mem[8] | | |
| 0 | | hit | Mem[0] | Mem[8] | | |
| 6 | | miss | Mem[0] | Mem[8] | Mem[6] | |
| 8 | | hit | Mem[0] | Mem[8] | Mem[6] | |

How Much Associativity

- Simulation of a system with 64KB D-cache, 16-word blocks, SPEC2000: (data miss rate)
 - 1-way: 10.3%
 - 2-way: 8.6%
 - 4-way: 8.3%
 - 8-way: 8.1%

Set Associative Cache Organization



Associativity Implies Choices

- Direct-mapped
 - Compare addr. with only one tag
 - Location A can be stored in exactly one cache line
- N-way set-associative
 - Compare addr. with N tags simultaneously
 - Location A can be stored in exactly one set, but in any of the N cache lines belonging to that set
- Fully associative
 - Compare addr. With each tag simultaneously
 - Location A can be stored in any cache line

Replacement Policy

- Direct mapped: no choice
- Set associative
 - Prefer non-valid entry, if there is one
 - Otherwise, choose among entries in the set
- Optimal policy: Replace the block that is accessed furthest in the future
 - Requires knowing the future
- Predict the future from looking at the past
 - If a block has not been used recently, it's often less likely to be accessed in the near future (a locality argument)
- Least-recently used (LRU)
 - Choose the one unused for the longest time

Write Policy

- Write-through:
 - CPU writes are cached, but also written to main memory immediately
 - Stalling the CPU until write is completed
 - Simple, slow
- Write-back:
 - CPU writes are cached, but not written to main memory until we replace the block
 - Commonly implemented in current systems
 - Fast, more complex

```
sum = 0;
for (i = 0; i < n; i++)
    sum += a[i];
return sum;
```

Write-Back with 'Dirty' Bits

- Add 1 bit per block to record whether block has been written to.
- Only write back dirty blocks