

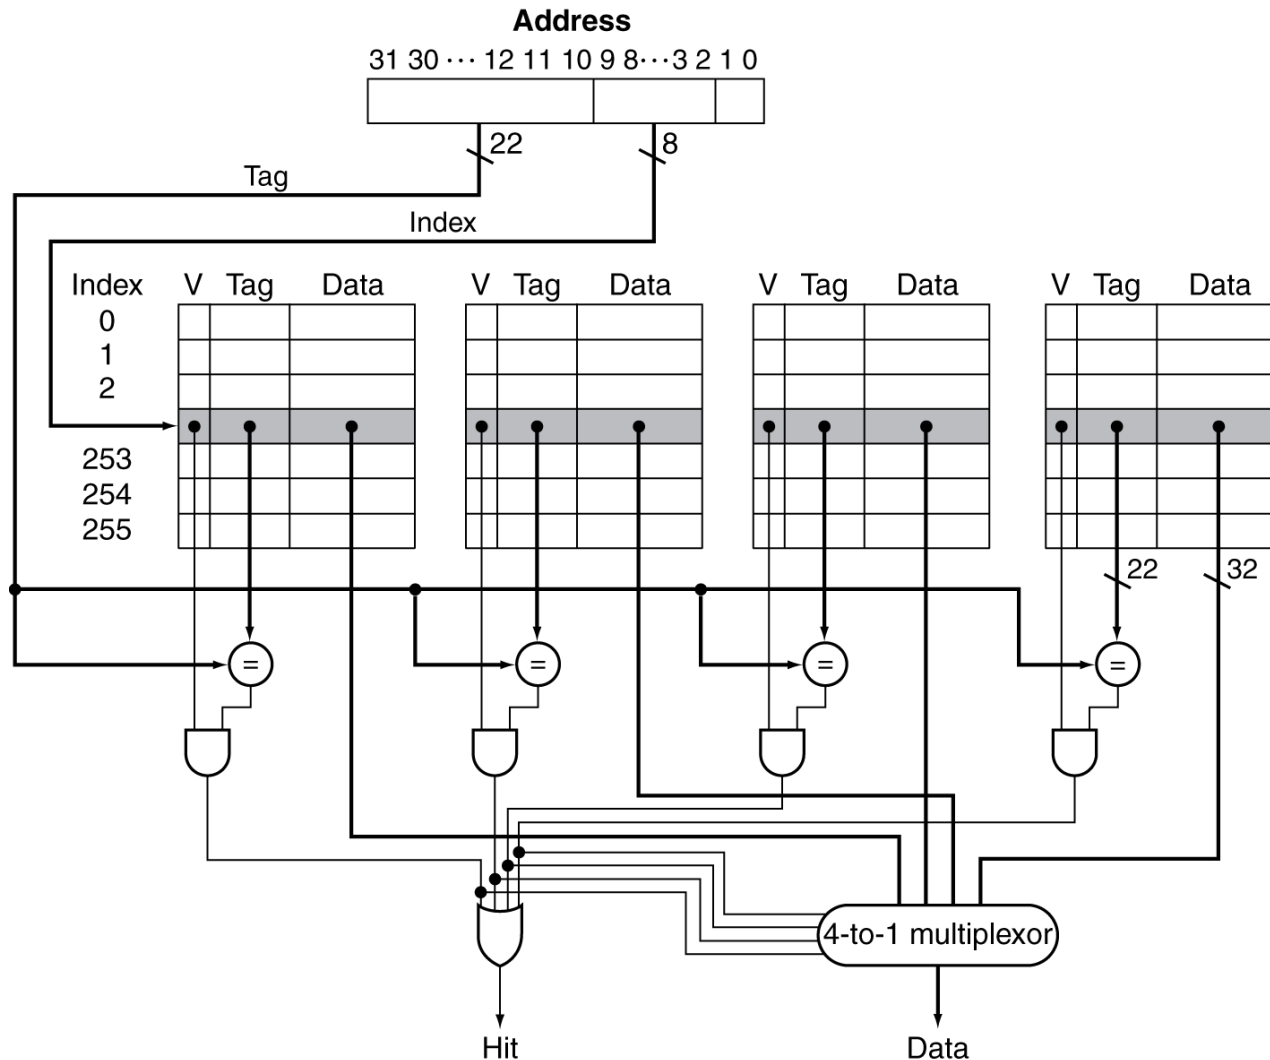


Memory Management

How Much Associativity

- Simulation of a system with 64KB D-cache, 16-word blocks, SPEC2000: (data miss rate)
 - 1-way: 10.3%
 - 2-way: 8.6%
 - 4-way: 8.3%
 - 8-way: 8.1%

Set Associative Cache Organization



Associativity Implies Choices

- Direct-mapped
 - Compare addr. with only one tag
 - Location A can be stored in exactly one cache line
- N-way set-associative
 - Compare addr. with N tags simultaneously
 - Location A can be stored in exactly one set, but in any of the N cache lines belonging to that set
- Fully associative
 - Compare addr. with each tag simultaneously
 - Location A can be stored in any cache line

Replacement Policy

- Direct mapped: no choice
- Set associative
 - Prefer non-valid entry, if there is one
 - Otherwise, choose among entries in the set
- Optimal policy: Replace the block that is accessed furthest in the future
 - Requires knowing the future
- Predict the future from looking at the past
 - If a block has not been used recently, it's often less likely to be accessed in the near future (a locality argument)
- Least-recently used (LRU)
 - Choose the one unused for the longest time

Write Policy

- Write-through:
 - CPU writes are cached, but also written to main memory immediately
 - Stalling the CPU until write is completed
 - Simple, slow
- Write-back:
 - CPU writes are cached, but not written to main memory until we replace the block
 - Commonly implemented in current systems
 - Fast, more complex

```
sum = 0;  
for (i = 0; i < n; i++)  
    sum += a[i];  
return sum;
```

Write-Back with 'Dirty' Bits

- Add 1 bit per block to record whether block has been written to.
- Only write back dirty blocks

Virtual Memory

- Software caches:
- Same objective: fake large, fast, and cheap memory
- Conceptually similar
- Different implementations



Virtual Memory

- Use main memory as a “cache” for secondary (disk) storage
 - Managed jointly by CPU hardware and the operating system (OS)
- Programs share main memory
 - Each gets a private virtual address space holding its frequently used code and data
 - Protected from other programs
- CPU and OS translate virtual addresses to physical addresses
 - VM “block” is called a page
 - VM translation “miss” is called a page fault

Virtual Memory

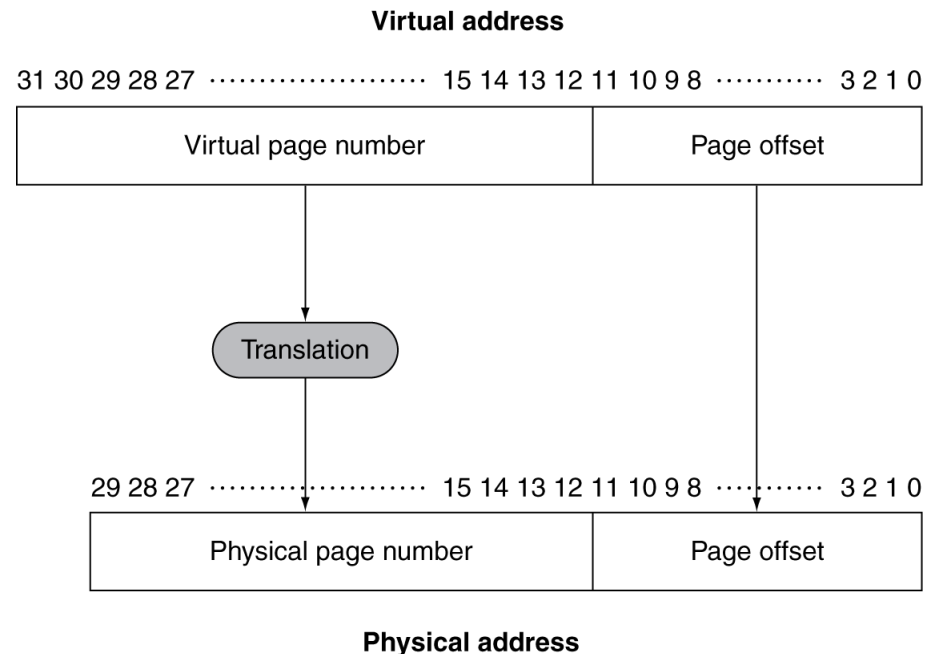

- Two kinds of addresses:
 - CPU (also programs) uses virtual addresses
 - Main memory uses physical addresses
- Hardware translates virtual addresses to physical addresses via an operating system(OS)-managed table, the **page map or page table**.
- The price of VM is address translation

Memory Management Unit



Address Translation

- Assume we have 1 GB main memory: how many bits required to represent physical address?
- Fixed-size pages (e.g., 4K): how many bits to describe page size?



Virtual Memory: Paging

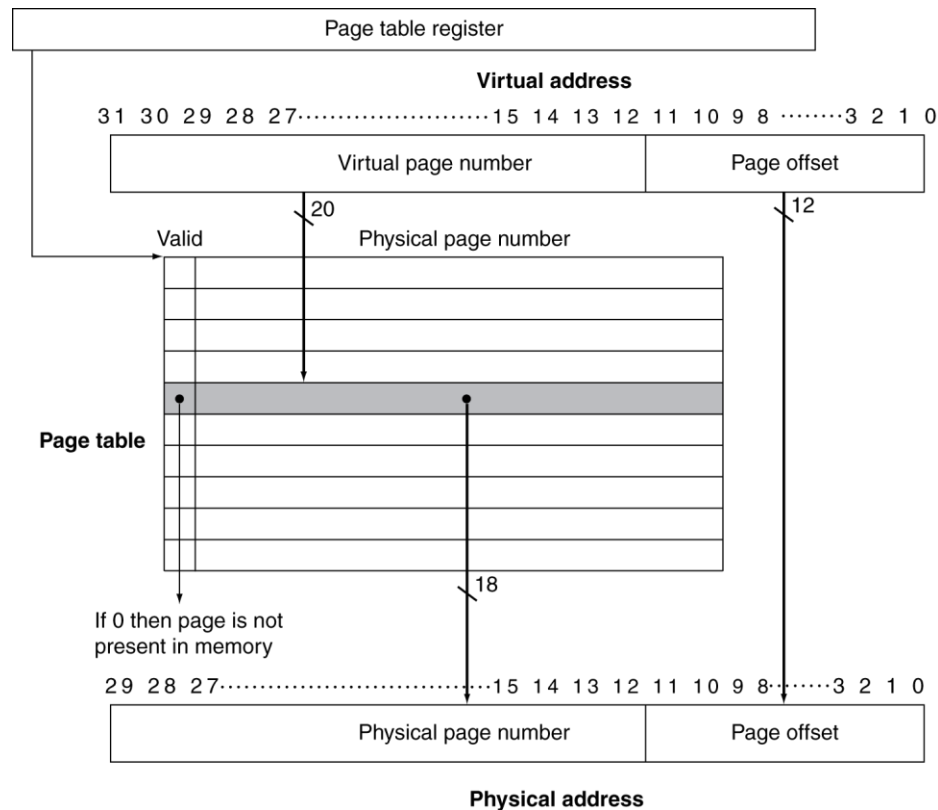
- Divide physical memory in fixed-sized blocks, called pages
 - Typical page size: 4~16KB
 - Virtual address: Virtual page number + offset bits
 - Physical address: Physical page number + offset bits
- MMU maps virtual address to physical address
 - Use page table (or page map) to perform translation
 - Miss: or called a **page fault** if no translation, which means the page is not present in main memory

Page Fault Penalty

- On page fault, the page must be fetched from disk
 - Takes millions of clock cycles
 - Handled by OS code
- Try to minimize page fault rate
 - Fully associative placement
 - Smart replacement algorithms

Page Tables

- Stores placement information
 - An array of page table entries(PTE), indexed by virtual page number
 - Page table register in CPU points to page table

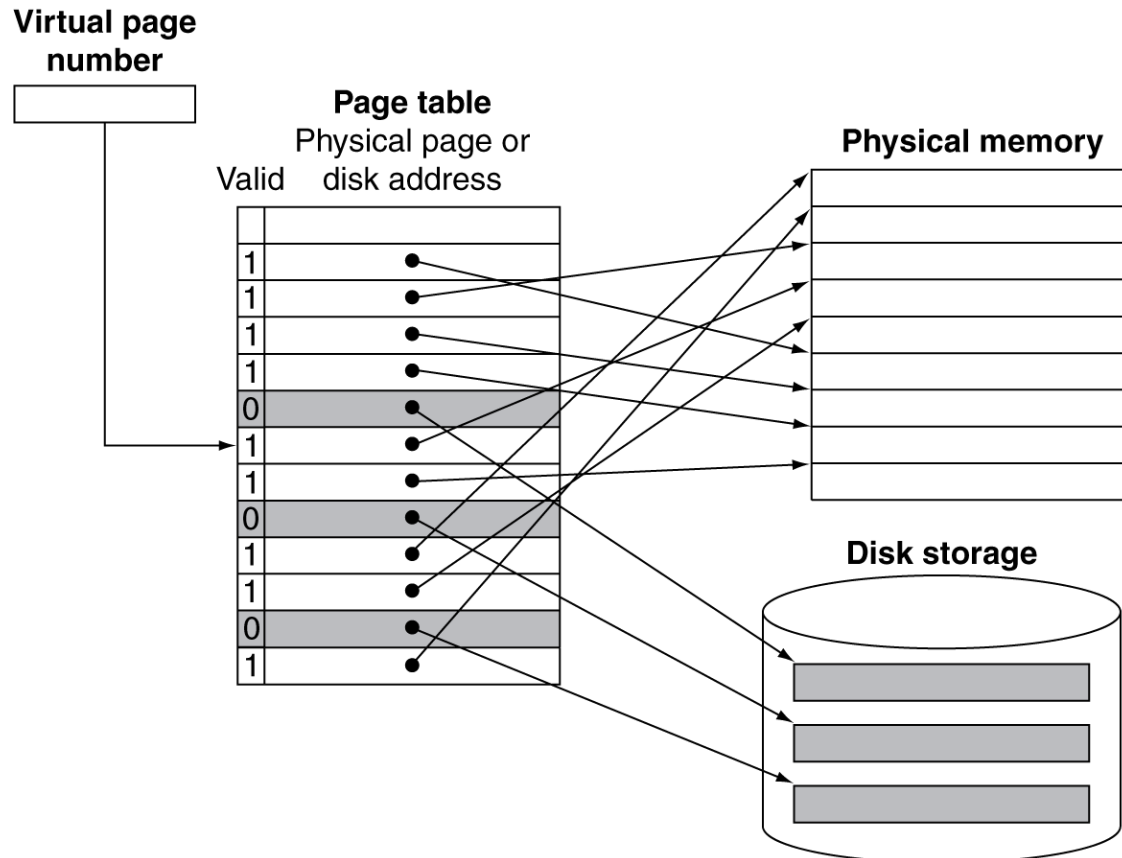


Page Tables

- If page is present in memory
 - PTE(Page Table Entry) stores the physical page number
 - Plus other status bits (valid, ...)
- If page is not present in memory
 - look for the page on disk

Page Tables

- Virtual page number serves as the index
- What if page fault and there is no empty page in physical memory.

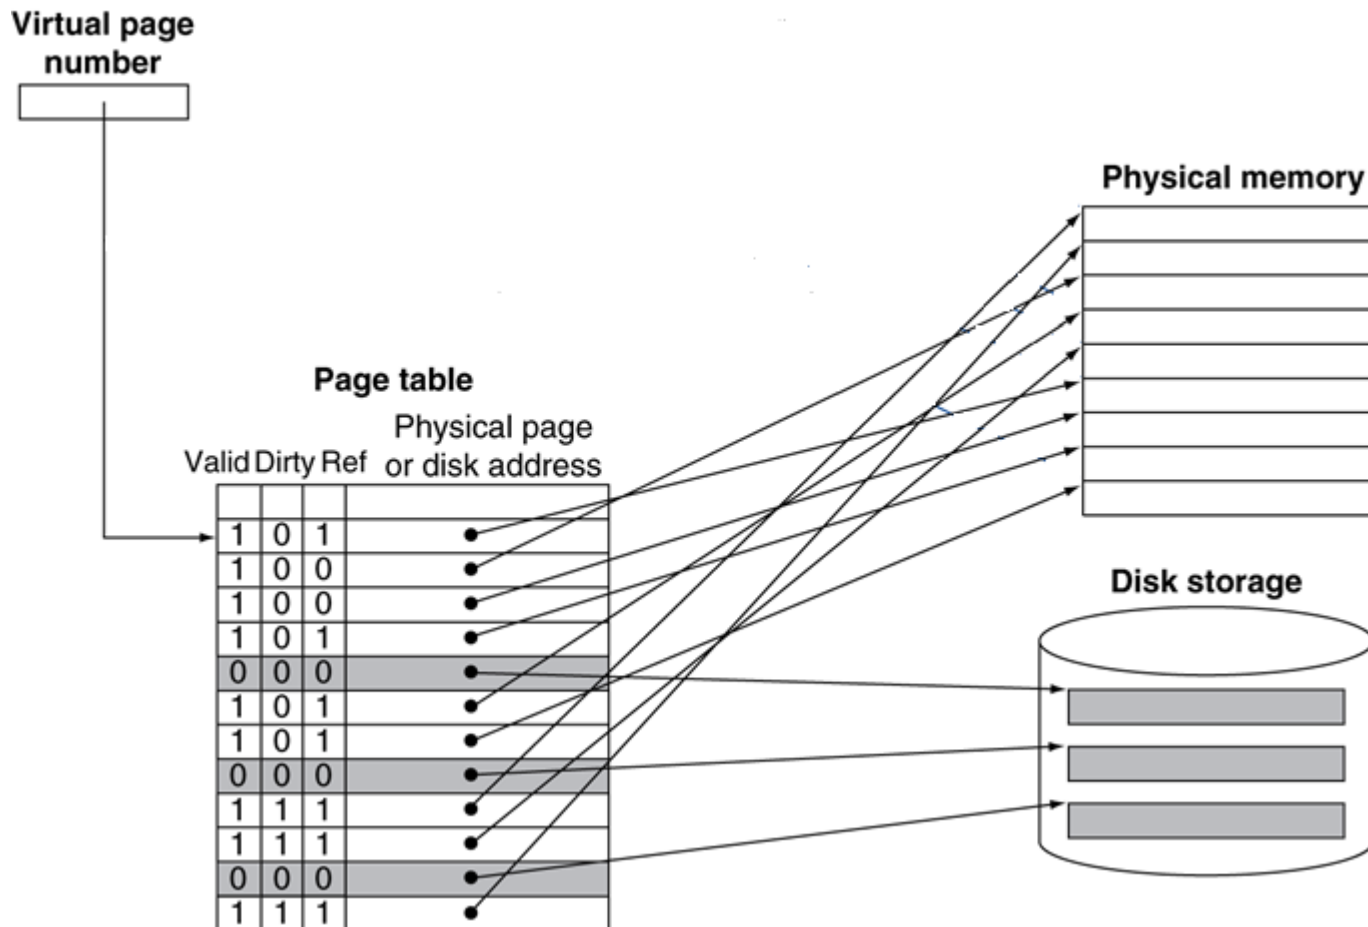


Replacement and Writes

- Replacement: To reduce page fault rate, prefer least-recently used (LRU) replacement
 - Reference bit (aka use bit) in PTE set to 1 on access to page
 - Periodically cleared to 0 by OS
 - A page with reference bit = 0 has not been used recently
- Disk writes take millions of cycles
 - Use write-back
 - Dirty bit in PTE set when page is written
- Two more bits for each entry in the page table

Page Tables

- Physical address $m + p$: m #bits represent page #
- How many bits in the page table:



Page Fault Handler

- Use virtual address to find PTE
- Locate page on disk
- Choose a page in memory to replace
 - If dirty, write to disk first
- Read page into memory and update page table
- Make process runnable again

Page Map Arithmetic

- $(v+p)$ bits in virtual address
- $(m+p)$ bits in physical address
- 2^v # of virtual pages
- 2^m # of physical pages
- 2^p bytes per physical page
- 2^{m+p} bytes in physical memory
- $(m+3) * 2^v$ bits in the page table:
- typical page size: 4~16KB
- typical $(v+p)$: 32 bit, 64 bit
- typical $(m+p)$: 30~40 bits (1GB~1TB)

Example: Page Map Arithmetic

- Suppose:
 - 32-bit virtual address
 - Page size: 4KB
 - 2^{12}
 - RAM max: 1GB
 - 2^{30}
- Then:
 - # Physical pages: 2^{18} 256K
 - # Virtual pages: 2^{20}
 - # PTE (page table entries): 2^{20}
 - # bits in page table: $(18+3)*2^{20}$
- Use SRAM for page table???

- 21Mbit → 3MB

